


KV cache: 12K tokens preserved



`read_file`

`edit_code`

`run_test`

`edit_code`

45 ms

180 ms

2.4 s

180 ms

bar width $\propto \sqrt{\text{idle time}}$

● *53× variance motivates TTL prediction*